

Exercise

For the dataset "multicollinearity.txt", perform the following:

1. Check the pairwise correlations of the predictors.
2. Compare the coefficient for each predictor in two cases: a. when regressing Y on that predictor only, and b. when regressing Y on all predictors. Do you see a major change in the coefficients? Their standard errors? What does this tell you about the presence of multicollinearity?
3. Calculate the VIFs for each predictor. Which VIFs are concerning? Comment on what you find.
4. Consider regressing `X4` onto the other predictors `X1` through `X3`. Save the residuals from this model and call them `res4`. Then regress Y onto `res4` and look at the coefficient. Compare this coefficient to the coefficients for `X4` you found in part 2. What do you notice?
5. Explain geometrically what we are doing in #4. Are we surprised by the results?

```
In [34]: data = pd.read_csv("~/Desktop/repos/regression_f24/data/multicollinearity.tx
```

```
In [35]: data.head()
```

```
Out[35]:
```

	y	x1	x2	x3	x4
0	1.860	0.374	0.846	1.951	1.325
1	9.747	2.184	3.754	5.099	5.372
2	12.707	2.164	3.635	2.505	2.582
3	12.822	5.595	4.805	3.653	3.399
4	5.160	5.330	2.611	1.873	1.474

```
In [36]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0    y      200 non-null      float64
1   x1      200 non-null      float64
2   x2      200 non-null      float64
3   x3      200 non-null      float64
4   x4      200 non-null      float64
dtypes: float64(5)
memory usage: 7.9 KB

```

```
In [37]: data.iloc[:, 1:].corr()
```

```

Out[37]:

```

	x1	x2	x3	x4
x1	1.000000	0.999811	0.999495	0.998843
x2	0.999811	1.000000	0.999671	0.999013
x3	0.999495	0.999671	1.000000	0.999306
x4	0.998843	0.999013	0.999306	1.000000

The coefficient for X1 goes from being positive at 2.25 to being negative at -0.1675 when regressed alongside all predictors. The std error grows by a factor of ~ 50x when other correlated predictors are added to the regression alongside x1.

```

In [39]: mod_y_x1 = ols("y ~ x1", data=data).fit()
mod_y_x1.summary()

```

Out [39]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999
Model:	OLS	Adj. R-squared:	0.999
Method:	Least Squares	F-statistic:	1.396e+05
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	5.39e-284
Time:	22:02:42	Log-Likelihood:	-589.73
No. Observations:	200	AIC:	1183.
Df Residuals:	198	BIC:	1190.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.7257	0.659	2.617	0.010	0.426	3.026
x1	2.1250	0.006	373.575	0.000	2.114	2.136

Omnibus:	0.666	Durbin-Watson:	2.009
Prob(Omnibus):	0.717	Jarque-Bera (JB):	0.451
Skew:	-0.106	Prob(JB):	0.798
Kurtosis:	3.095	Cond. No.	233.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [40]: mod_y_all = ols("y ~ x1 + x2 + x3 + x4", data=data).fit()
mod_y_all.summary()
```

Out [40]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999
Model:	OLS	Adj. R-squared:	0.999
Method:	Least Squares	F-statistic:	7.200e+04
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	8.92e-308
Time:	22:02:42	Log-Likelihood:	-515.82
No. Observations:	200	AIC:	1042.
Df Residuals:	195	BIC:	1058.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.0374	0.462	2.248	0.026	0.127	1.948
x1	-0.1675	0.204	-0.821	0.413	-0.570	0.235
x2	1.1364	0.281	4.050	0.000	0.583	1.690
x3	0.9788	0.256	3.830	0.000	0.475	1.483
x4	1.1175	0.210	5.311	0.000	0.703	1.533

Omnibus:	0.460	Durbin-Watson:	2.103
Prob(Omnibus):	0.795	Jarque-Bera (JB):	0.219
Skew:	-0.046	Prob(JB):	0.896
Kurtosis:	3.133	Cond. No.	377.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The coefficient for X2 changes but stays positive. The standard error increases by a factor of 60 times (!). This is variance inflation!

```
In [42]: mod_y_x2 = ols("y ~ x2", data=data).fit()
mod_y_x2.summary()
```

Out [42]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999			
Model:	OLS	Adj. R-squared:	0.999			
Method:	Least Squares	F-statistic:	1.980e+05			
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	5.20e-299			
Time:	22:02:42	Log-Likelihood:	-554.80			
No. Observations:	200	AIC:	1114.			
Df Residuals:	198	BIC:	1120.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.3706	0.554	2.473	0.014	0.278	2.464
x2	2.3613	0.005	444.944	0.000	2.351	2.372
Omnibus:	0.045	Durbin-Watson:	1.978			
Prob(Omnibus):	0.978	Jarque-Bera (JB):	0.043			
Skew:	0.030	Prob(JB):	0.979			
Kurtosis:	2.960	Cond. No.	210.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The coefficient changes, but sign stays the same. The SE grows by a factor of 50x...!

```
In [44]: mod_y_x3 = ols("y ~ x3", data=data).fit()
mod_y_x3.summary()
```

Out [44]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999			
Model:	OLS	Adj. R-squared:	0.999			
Method:	Least Squares	F-statistic:	2.218e+05			
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	6.86e-304			
Time:	22:02:42	Log-Likelihood:	-543.45			
No. Observations:	200	AIC:	1091.			
Df Residuals:	198	BIC:	1098.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.0831	0.524	2.066	0.040	0.049	2.117
x3	2.9492	0.006	470.945	0.000	2.937	2.962
Omnibus:	0.792	Durbin-Watson:	2.128			
Prob(Omnibus):	0.673	Jarque-Bera (JB):	0.741			
Skew:	-0.148	Prob(JB):	0.690			
Kurtosis:	2.965	Cond. No.	169.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [45]: m_y_x4 = ols("y ~ x4", data=data).fit()
m_y_x4.summary()
```

Out [45]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.998
Model:	OLS	Adj. R-squared:	0.998
Method:	Least Squares	F-statistic:	1.307e+05
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	3.46e-281
Time:	22:02:42	Log-Likelihood:	-596.26
No. Observations:	200	AIC:	1197.
Df Residuals:	198	BIC:	1203.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.8964	0.683	1.312	0.191	-0.451	2.244
x4	4.2002	0.012	361.558	0.000	4.177	4.223

Omnibus:	0.500	Durbin-Watson:	2.234
Prob(Omnibus):	0.779	Jarque-Bera (JB):	0.641
Skew:	0.079	Prob(JB):	0.726
Kurtosis:	2.772	Cond. No.	119.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

VIF

The VIF values for all the predictors are very high, pointing to severe multicollinearity in the model, confirming the conclusion of the previous question

```
In [48]: def calculate_VIF(r2_j):
         return (1/(1-r2_j))
```

```
In [49]: mdl_x1 = ols("x1~x2+x3+x4", data=data).fit()
         r2j_x1 = mdl_x1.rsquared
         print(r2j_x1)
         print(calculate_VIF(r2j_x1))
```

```
0.9996231915946605
2653.8686128804293
```

```
In [50]: mdl_x2 = ols("x2~x1+x3+x4", data=data).fit()
r2j_x2 = mdl_x2.rsquared
print(r2j_x2)
print(calculate_VIF(r2j_x2))
```

```
0.9997540700184278
4066.1980032159418
```

```
In [51]: mdl_x3 = ols("x3~x2+x1+x4", data=data).fit()
r2j_x3 = mdl_x3.rsquared
print(r2j_x3)
print(calculate_VIF(r2j_x3))
```

```
0.9995375633481587
2162.4583519023327
```

```
In [52]: mdl_x4 = ols("x4~x2+x3+x1", data=data).fit()
r2j_x4 = mdl_x4.rsquared
print(r2j_x4)
print(calculate_VIF(r2j_x4))
```

```
0.9986155124710678
722.2889185367274
```


4. $res4 = X_4 - \hat{X}_4$ $X_4 \sim X_1 + X_2 + X_3$

The coefficient for $res4$ is the same as that of X_4 in the model with all the predictors and of $res4$ in the model with all the other predictors. (Key observation: $\underline{res4} \perp \text{span}\{X_1, X_2, X_3\}$.)

```
In [54]: modres = ols("x4~x1+x2+x3", data=data).fit()
res4 = modres.resid
```

```
In [55]: data['res4'] = res4
```

```
In [56]: mod_y_res4 = ols("y~res4", data=data).fit()
mod_y_res4.summary()
```



Out [56]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.000
Model:	OLS	Adj. R-squared:	-0.005
Method:	Least Squares	F-statistic:	0.01938
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	0.889
Time:	22:02:42	Log-Likelihood:	-1245.7
No. Observations:	200	AIC:	2495.
Df Residuals:	198	BIC:	2502.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	215.3599	8.716	24.707	0.000	198.171	232.549
res4	1.1175	8.028	0.139	0.889	-14.714	16.949
Omnibus:	91.438	Durbin-Watson:	0.003			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	11.967			
Skew:	-0.004	Prob(JB):	0.00252			
Kurtosis:	1.802	Cond. No.	1.09			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [57]: `mod_y_x123res4 = ols("y~x1+x2+x3+res4", data=data).fit()
mod_y_x123res4.summary()`

Out [57]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999
Model:	OLS	Adj. R-squared:	0.999
Method:	Least Squares	F-statistic:	7.200e+04
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	8.92e-308
Time:	22:02:42	Log-Likelihood:	-515.82
No. Observations:	200	AIC:	1042.
Df Residuals:	195	BIC:	1058.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.1484	0.461	2.491	0.014	0.239	2.058
x1	-0.1589	0.204	-0.779	0.437	-0.561	0.244
x2	1.1616	0.281	4.140	0.000	0.608	1.715
x3	1.7193	0.214	8.027	0.000	1.297	2.142
res4	1.1175	0.210	5.311	0.000	0.703	1.533

Omnibus:	0.460	Durbin-Watson:	2.103
Prob(Omnibus):	0.795	Jarque-Bera (JB):	0.219
Skew:	-0.046	Prob(JB):	0.896
Kurtosis:	3.133	Cond. No.	358.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [58]: mod_y_all = ols("y ~ x1 + x2 + x3 + x4", data=data).fit()
mod_y_all.summary()
```

Out [58]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.999
Model:	OLS	Adj. R-squared:	0.999
Method:	Least Squares	F-statistic:	7.200e+04
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	8.92e-308
Time:	22:02:42	Log-Likelihood:	-515.82
No. Observations:	200	AIC:	1042.
Df Residuals:	195	BIC:	1058.
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.0374	0.462	2.248	0.026	0.127	1.948
x1	-0.1675	0.204	-0.821	0.413	-0.570	0.235
x2	1.1364	0.281	4.050	0.000	0.583	1.690
x3	0.9788	0.256	3.830	0.000	0.475	1.483
x4	1.1175	0.210	5.311	0.000	0.703	1.533

Omnibus:	0.460	Durbin-Watson:	2.103
Prob(Omnibus):	0.795	Jarque-Bera (JB):	0.219
Skew:	-0.046	Prob(JB):	0.896
Kurtosis:	3.133	Cond. No.	377.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

But note the coefficient is different from that in the model with just y onto x4:

In [60]: `m_y_x4.summary()`

$$y \sim x_4$$

Out [60]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.998
Model:	OLS	Adj. R-squared:	0.998
Method:	Least Squares	F-statistic:	1.307e+05
Date:	Wed, 25 Sep 2024	Prob (F-statistic):	3.46e-281
Time:	22:02:42	Log-Likelihood:	-596.26
No. Observations:	200	AIC:	1197.
Df Residuals:	198	BIC:	1203.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.8964	0.683	1.312	0.191	-0.451	2.244
x4	4.2002	0.012	361.558	0.000	4.177	4.223

Omnibus:	0.500	Durbin-Watson:	2.234
Prob(Omnibus):	0.779	Jarque-Bera (JB):	0.641
Skew:	0.079	Prob(JB):	0.726
Kurtosis:	2.772	Cond. No.	119.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

5.

Geometric Interpretation of the Models

Model 1:

$$res4 = \tilde{x}_4$$

$$y \sim X_1 + X_2 + X_3 + res4:$$

Here, $res4$ is the vector that represents the part of

X_4 that is orthogonal to $\text{span}\{X_1, X_2, X_3\}$. Since it is orthogonal to $\text{span}\{X_1, X_2, X_3\}$, it adds information to the model that $\text{span}\{X_1, X_2, X_3\}$ cannot explain. Geometrically, you are using four vectors: X_1 , X_2 , X_3 , and $res4$, form a 4-dimensional basis.

Model 2:

$$y \sim X_1 + X_2 + X_3 + X_4:$$

In this model, you are using X_4 directly. However, when you regress y onto all predictors X_1, X_2, X_3, X_4 , any part of

X_4 that is explained by X_1, X_2 , or X_3 (the projection of X_2 onto $\text{span}\{X_1, X_2, X_3\}$) does not add new information. The model is effectively using the part of X_4 which is orthogonal to $\text{span}\{X_1, X_2, X_3\}$ (i.e., the same res_4 from the first model) to explain y , in addition to X_1, X_2 , and X_3 .

$$\text{res}_2 = X_2 - \hat{X}_2 = X_2 - \tilde{\beta}_0 - \tilde{\beta}_1 X_1$$

In both models, the last predictor variable (whether it's res_4 in Model 1 or X_4 in Model 2) represents the same direction in the vector space: the part of X_4 that is orthogonal to $\text{span}\{X_1, X_2, X_3\}$

In $y \sim X_1 + X_2 + X_3 + \text{res}_4$, you're explicitly using the orthogonal part of X_4 (i.e., res_4). In $y \sim X_1 + X_2 + X_3 + X_4$, although you include all of X_4 , the model implicitly separates the part of X_4 that is redundant (explained by $\{X_1, X_2, X_3\}$) and only uses the part of X_4 orthogonal to $\text{span}\{X_1, X_2, X_3\}$, just like in the first model.

$$y \sim X_1 + X_2$$

$$y \sim X_1 + \text{res}_2$$

$$H = (I \ X_1 \ X_2) \quad HY = \hat{y}$$

In []:

$\text{span}(X_1, \text{res}_2)$
 $\text{span}(X_1, X_2)$

$\text{Colspace}(X)$

Full

$$y \sim X_1 + X_2$$

$$X_2 \sim X_1$$

$$\hookrightarrow \text{res}_2$$

$$y \sim X_1 + \text{res}_2$$

Reduced:

$$y \sim X_1$$

$$(I-H)y \quad (I-H)X_2$$

$$\langle e_f, \text{res}_2 \rangle =$$

$$y(I-H)^T(I-H)X_2 = \dots = 0$$

$$\|res_2\|^2 + \|e_f\|^2 = \|e_r\|^2 \leftarrow$$

$$\|res_2\|^2 + SSE_f = SSE_r$$

$$SS(x_2|x_1) + \|res_2\|^2 = SSE_r - SSE_f$$

$$F = \frac{SS(x_2|x_1)/1}{\hat{\sigma}^2} \sim t^2$$

$$* \quad res_2 = X_2 - H_{(x_1)} X_2 = (I - H_{(x_1)}) X_2$$

$$e_f = Y - HY = (I - H)Y$$

$$\begin{aligned} \langle e_f, res_2 \rangle &= Y(I - H)(I - H_{(x_1)}) X_2 \\ &= Y(I - H - H_{(x_1)} + HH_{(x_1)}) X_2 \\ &= Y(I - H - H_{(x_1)} + H_{(x_1)}) X_2 = Y(I - H) X_2 = Y(X_2 - HX_2) = Y(X_2 - X_2) = 0 \end{aligned}$$

why is $H_{(x_1)} = H_{(x_1)}$?


$$\begin{aligned} H_{(x_1)} &= HX_1(X_1^T X_1)^{-1} X_1^T \\ &= X_1(X_1^T X_1)^{-1} X_1^T = H_{(x_1)} \end{aligned} \quad \begin{aligned} X_1 &\in \text{span}(X_1, X_2) \\ &\Rightarrow HX_1 = X_1 \end{aligned}$$

If $X_1 \perp X_2$,

$$X_2 \sim X_1 \Rightarrow$$

$$res_2 = X_2 - \bar{X}_2$$


= X_2 shifted by a constant



Understanding Multicollinearity with Spgs


If $X_1 \perp X_2$, then

$Y \sim X_1 + X_2$ looks like:




if X_1 & X_2 are very collinear,
then $y \sim X_1 + X_2$ looks

like:




many planes result in similar
SSEs! $\Rightarrow$ small changes in X result
in large changes in estimated
slopes $\hat{\beta}$!



$$Full: y \sim x_1 + x_2$$

$$\hat{y}_H = H_F y$$



$$\hat{x}_3 = x_3$$

$$\hat{x}_3 \sim x_1 + x_2$$

$$\text{span}(x_1, x_2)$$

$$\|e_f\|^2 = \underbrace{\|e_{sf}\|^2} + \|res_3\|^2$$

$$= \|e_r\|^2 + \underbrace{\|res_2\|^2} + \|res_3\|^2$$